

The Adventures of Malloc and New

Lecture 3: Oh, Say Can You C

Eunsuk Kang and **Jean Yang**

MIT CSAIL

January 21, 2010

Homework notes



Courtesy of xkcd.com. Comic is available here: <http://xkcd.com/336/>

Programs must:

- Compile and run (with instructions) to receive a ✓.
- Compile and run, passing tests, to receive a ✓+.

Lecture plan

The more I C, the less I see. –Unknown.

1. Review of main concepts from previous lectures.
2. Fancier memory examples.
3. Closer look at GCC.
4. Style and tips.
5. Why C?

Review: stack and heap

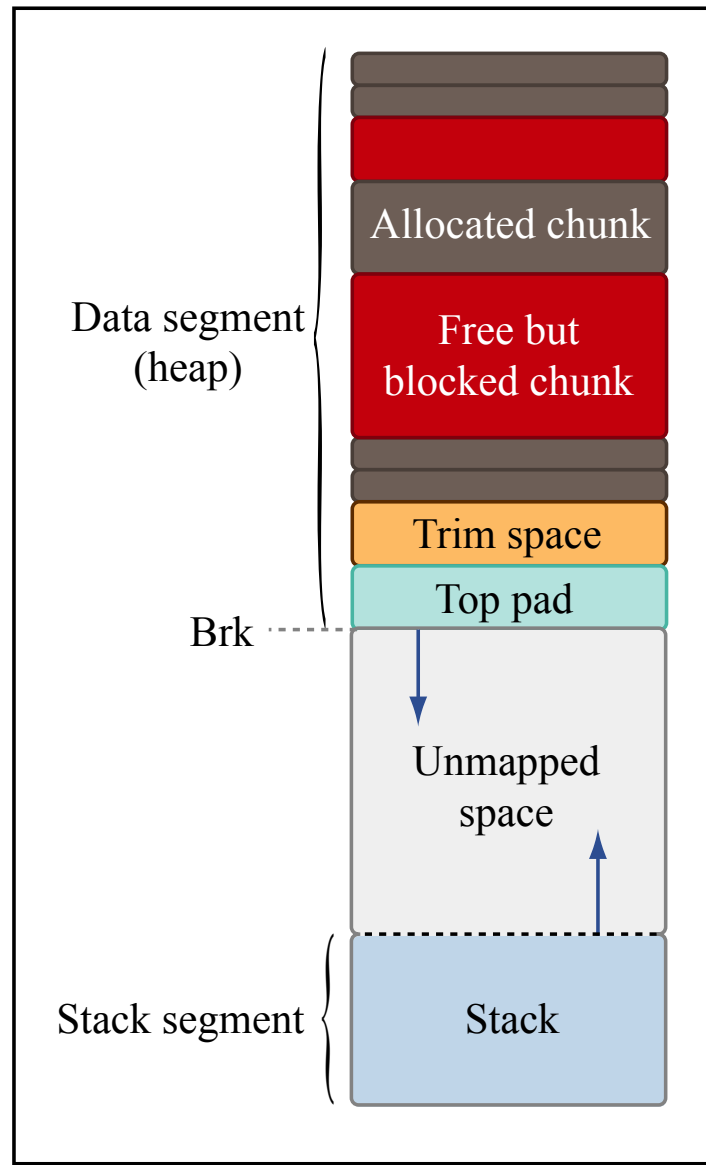


Figure by MIT OpenCourseWare.

Review: when to use pointers

- When you have to allocate memory on heap. (When is this?)
- When passing a parameter whose value you want to allow the other function to change.
- Also for efficiency—to avoid copying data structures.

Buggy field access

```
int* i = NULL;  
*i = 3;
```

Segmentation fault!

```
struct pair {  
    int first;  
    int second;  
};  
  
struct pair* pp = NULL;  
pp->first = 1;
```

Segmentation fault!

Buggy free

```
struct pair* pp = malloc(sizeof(struct pair));  
pp = NULL;  
free(pp);
```

Memory leak!

```
int* i = NULL;  
free(i);
```

Nothing bad happens! Freeing NULL does nothing.

Buggy scope

Buggy example

```
int* get_array(int* len) {  
    *len = 3;  
    return vals;  
}  
  
int main() {  
    int len, i;  
    int* arr = get_array(&len);  
    for (i = 0; i < len; +i) {  
        printf("%d\n", arr[i]);  
    }  
    return 0;  
}
```

Returns address of local (statically allocated) variable. (Should get warning!)

Buggy scope continued...

Buggy output

1
32607
-1527611392

Correct program

```
int* get_array(int* len) {  
    *len = 3;  
    int* vals = malloc(sizeof(int) * 3);  
    arr[0] = 1; arr[1] = 2; arr[2] = 3;  
    return arr;  
}
```

Buggy initialization

```
struct pair* pp;  
int i = pp->first;
```

Maybe a segmentation fault...

In-place linked list reversal

```
Element *reverse(Element *old_list)
{
    Element *new_list = NULL;

    while (old_list != NULL) {
        // Remove element from old list.
        Element *element = old_list;
        old_list = old_list->next;

        // Insert element in new list.
        element->next = new_list;
        new_list = element;
    }

    return new_list;
}
```

Courtesy of Lawrence Kesteloot. Used with permission.

Please see http://www.teamten.com/lawrence/writings/reverse_a_linked_list.html.

Constant time insert into a circular singly-linked list




Figure by MIT OpenCourseWare.

- Circular linked list: last node has a pointer to the first node.
- Given a pointer to a node—can't change that pointer!
- Want to insert a node *before* the current one—can we do that in constant time?

A closer look at the GCC compilation process

Preprocessor

Translation of `#` directives.

- Translates all macros (`#DEFINE`'s) into inline C code.
- Takes `#include` files and inserts them into the code.
 - Get redefinition error if structs etc. are defined more than once!
 - Use `#ifndef` directive to define things only if they have not been defined.

```
#ifndef __HEADER_NAME
#define __HEADER_NAME
/* Header code here. */
#endif
```

Aside: #define

Compiler directive.

```
#define DEFINED_CONSTANT 3
```

```
#define increment(x) (x+1)
```

GCC continued

Parsing and translation

Translates to assembly, performing optimizations.

Assembler

Translates assembly to machine instructions.

Linking

- **Static.** For each function called by the program, the assembly to that function is included directly in the executable, allowing function calls to directly address code.
- **Dynamic.** Function calls call a Procedure Linkage Table, which contains the proper addresses of the mapped memory.

Some helpful compiler options

Enforcements and warnings	
-ansi	Tells compiler to enforce ANSI C standards.
-pedantic	More pedantic ANSI with warnings.
-Wall	Shows all warnings.
-W	Show some warnings (return without a value, etc.).
Compilation/output	
-c	Compile and assemble, but do not link.
-S	Stop after compilation; do not assemble.
-E	Stop after preprocessing; do not compile.
-o [file]	Put the output binary in [file].
Profiling	
-pg	Compile with profiling information.


```
#define _ F-->00 || F--OO--;
long F=00,OO=00;
main(){F_OO();printf("%1.3f\n", 4.*-F/OO/OO);}F_OO()
{
    _ _ _ _ _
  _ _ _ _ _ _ _ _ _
 _ _ _ _ _ _ _ _ _ _ _
_ _ _ _ _ _ _ _ _ _ _ _ _
_ _ _ _ _ _ _ _ _ _ _ _ _
_ _ _ _ _ _ _ _ _ _ _ _ _
_ _ _ _ _ _ _ _ _ _ _ _ _
_ _ _ _ _ _ _ _ _ _ _ _ _
_ _ _ _ _ _ _ _ _ _ _ _ _
_ _ _ _ _ _ _ _ _ _ _ _ _
_ _ _ _ _ _ _ _ _ _ _ _ _
_ _ _ _ _ _ _ _ _ _ _ _ _
_ _ _ _ _ _ _ _ _ _ _ _ _
  _ _ _ _ _ _ _ _ _ _ _
    _ _ _ _ _ _ _ _ _ _
      _ _ _ _ _ _ _ _ _
        _ _ _ _ _ _ _
          _ _ _ _ _
            _ _ _
              _
                }

```

Courtesy of Landon Curt Noll and Larry Bassel. Used with permission.

Permission for personal, educational or non-profit use is granted provided this copyright and notice are included in its entirety and remains unaltered. All other uses must receive prior permission in writing from both Landon Curt Noll and Larry Bassel.

```
gcc pi.c -E
```

```
long F=00,OO=00;
main() {F_OO(); printf(" %1.3f\n" , 4.*-F/OO/OO);} F_OO()
{
    F-->00 || F_OO--;-F-->00 || F_OO--;-F-->00
        || F_OO--;-F-->00 || F_OO--;
    F-->00 || F_OO--;-F-->00 || F_OO--;-F-->00 || F-
        OO--;-F-->00 || F_OO--;-F-->00 || F_OO--;-F
        -->00 || F_OO--;-F-->00 || F_OO--;-F-->00 ||
        F_OO--;-F-->00 || F_OO--;
    F-->00 || F_OO--;-F-->00 || F_OO--;-F-->00 || F_OO
        --;-F-->00 || F_OO--;-F-->00 || F_OO--;-F-->00
        || F_OO--;-F-->00 || F_OO--;-F-->00 || F_OO--;-
        F-->00 || F_OO--;-F-->00 || F_OO--;-F-->00 || F
        _OO--;-F-->00 || F_OO--;
    F-->00 || F_OO--;-F-->00 || F_OO--;-F-->00 || F_OO
        --;-F-->00 || F_OO--;-F-->00 || F_OO--;-F-->00 ||
        F_OO--;-F-->00 || F_OO--;-F-->00 || F_OO--;-F
        -->00 || F_OO--;-F-->00 || F_OO--;-F-->00 || F_OO
        --;-F-->00 || F_OO--;-F-->00 || F_OO--;-F-->00 ||
        F_OO--;
```

Build systems: a quick flavor of make

- Way to manage compilation for large systems.
- Automatically determines which parts of large programs need to be recompiled.
- Simple makefile consists of rules of the following form:
target ... : prerequisites ...
 command
 ...
 ...
- Build system by running `make` from command line.

A simple C makefile

Makefile

This uses the implicit rules GNU Make defines for compiling C.

```
CC=gcc
CFLAGS=-Wall
main: main.o hello_fn.o
clean:
    rm -f main main.o hello_fn.o
```

Compiling with make

```
$ make
gcc -Wall -c -o main.o main.c
gcc -Wall -c -o hello_fn.o hello_fn.c
gcc main.o hello_fn.o -o main
$ ./main
"Hello world!"
```

Profiling: gprof

1. Compile with the profiling option.

```
gcc single_linked_list.c test_sll.c -o sll -pg
```

2. Run the program—this will produce a file gmon.out (unless otherwise specified) containing profiling information.

```
./sll
```

3. Run gprof on the binary to get the profiling information.
(Can suppress/select specific functions.)

```
gprof sll > profile_output.txt
```

Sample profiling output

Flat profile

% time	cumulative seconds	self seconds	calls	self Ts/call	total Ts/call	name
0.00	0.00	0.00	6	0.00	0.00	print_sll
0.00	0.00	0.00	5	0.00	0.00	make_node
0.00	0.00	0.00	4	0.00	0.00	insert_val
0.00	0.00	0.00	2	0.00	0.00	delete_val

Call tree

Also shows more detailed call tree information, sorted by total amount of time spent in each function and its children.

War stories




Figure: With great power comes great responsibility.

Courtesy of xkcd.com. Comic is available here: <http://xkcd.com/229/>

- Linker woes—why you want namespaces.
- You can link anything!

Write principled code in C

Cartoon of lemming falling off a cliff removed due to copyright restrictions.

I will not be a lemming and follow the crowd over the cliff and into the C. –John (Jack) Beidler

Array indexing

Recall that an array of size n with objects of type t is just a block of memory of size $\text{sizeof}(t) * n$.

Element	1	2	3	n
Array (<code>arr []</code>)	<code>arr[0]</code>	<code>arr[1]</code>	<code>arr[2]</code>	<code>arr[$n - 1$]</code>
Pointer (<code>arr*</code>)	<code>*arr</code>	<code>*(arr+1)</code>	<code>*(arr+2)</code>	<code>*(arr + $n + 1$)</code>

Style guidelines

- Test for equality with the constant on the left-hand side.

```
if (3 == x) /* rather than (x == 3) {  
    ...  
}
```

- Initialize pointers to NULL.

```
int* p = NULL;  
int* q; /* Unitialized , q will point to junk. */
```

- Use pre-increment unless post-increment is necessary.

```
++i; /* Pre-increment. */  
i++; /* Post-increment. */
```

Where is the overhead?

In memory-managed languages

- **Garbage collection**—figuring out what can be garbage collected and reclaiming that memory.
- Overhead from GC's *conservative* approximations of what is still in use.
- Reducing object allocation mitigates this problem.

In C

- Each **allocation** takes time because the allocator has to search for a piece of sufficiently large memory each time.
- Reducing the number of allocations mitigates this problem.

C is not just for thrill-seekers

Speed

- Slowdowns from compiling from higher-level languages.
- Personal anecdote: randomized simulations (C vs. Python).

Memory

- Memory overhead (20%+) necessary for garbage collection.
- Personal anecdote: 400x speedup from reducing object creation in OCaml program.

Access to dangerous, low-level parts of system

- Accessing hardware devices as memory.
- Mucking with the registers (stack pointer).

Until tomorrow...

Homework (due Saturday)

- Rewrite the binary search tree assignment using an array.
- Please submit a .zip or .tar.gz file prefaced by [username].hw[number].
- More details on the course website.

Questions?

- The course staff will be available after class.

MIT OpenCourseWare
<http://ocw.mit.edu>

6.088 Introduction to C Memory Management and C++ Object-Oriented Programming

January IAP 2010

For information about citing these materials or our Terms of Use, visit: <http://ocw.mit.edu/terms>.